

```
In [2]: ▶ import pandas as pd
import numpy as np
import plotly.express as px
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
```

## Reading external dataset

```
In [3]: ▶ df=pd.read_csv("salaries.csv")
```

```
In [4]: ▶ df.head()
```

Out[4]:

	work_year	experience_level	employment_type	job_title	salary	salary_currency	salary
0	2023	EX	FT	Data Science Director	212000	USD	
1	2023	EX	FT	Data Science Director	190000	USD	
2	2023	MI	FT	Business Intelligence Engineer	35000	GBP	
3	2023	MI	FT	Business Intelligence Engineer	35000	GBP	
4	2023	SE	FT	Machine Learning Engineer	245700	USD	

```
In [5]: ▶ df.shape
```

Out[5]: (8805, 11)

In [6]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 8805 entries, 0 to 8804
Data columns (total 11 columns):
#   Column                Non-Null Count  Dtype
---  -
0   work_year              8805 non-null   int64
1   experience_level        8805 non-null   object
2   employment_type        8805 non-null   object
3   job_title              8805 non-null   object
4   salary                 8805 non-null   int64
5   salary_currency        8805 non-null   object
6   salary_in_usd          8805 non-null   int64
7   employee_residence     8805 non-null   object
8   remote_ratio           8805 non-null   int64
9   company_location       8805 non-null   object
10  company_size            8805 non-null   object
dtypes: int64(4), object(7)
memory usage: 756.8+ KB
```

## Producing various descriptive statistics of the analytic dataset

In [7]: `df['salary_in_usd'].describe()`

```
Out[7]: count      8805.000000
mean      149488.265645
std       64222.105058
min       15000.000000
25%      105000.000000
50%      142200.000000
75%      185900.000000
max       615201.000000
Name: salary_in_usd, dtype: float64
```


```
In [8]: median_salary = df['salary_in_usd'].median()
Variance = df['salary_in_usd'].var()
salary_range = df['salary_in_usd'].max() - df['salary_in_usd'].min()
mode_salary = df['salary_in_usd'].mode()
```

```
In [9]: print("Median Salary: ", median_salary)
print("Salary Variance: ", Variance)
print("Salary Range: ", salary_range)
print("Mode Salary: ", mode_salary.tolist())
```

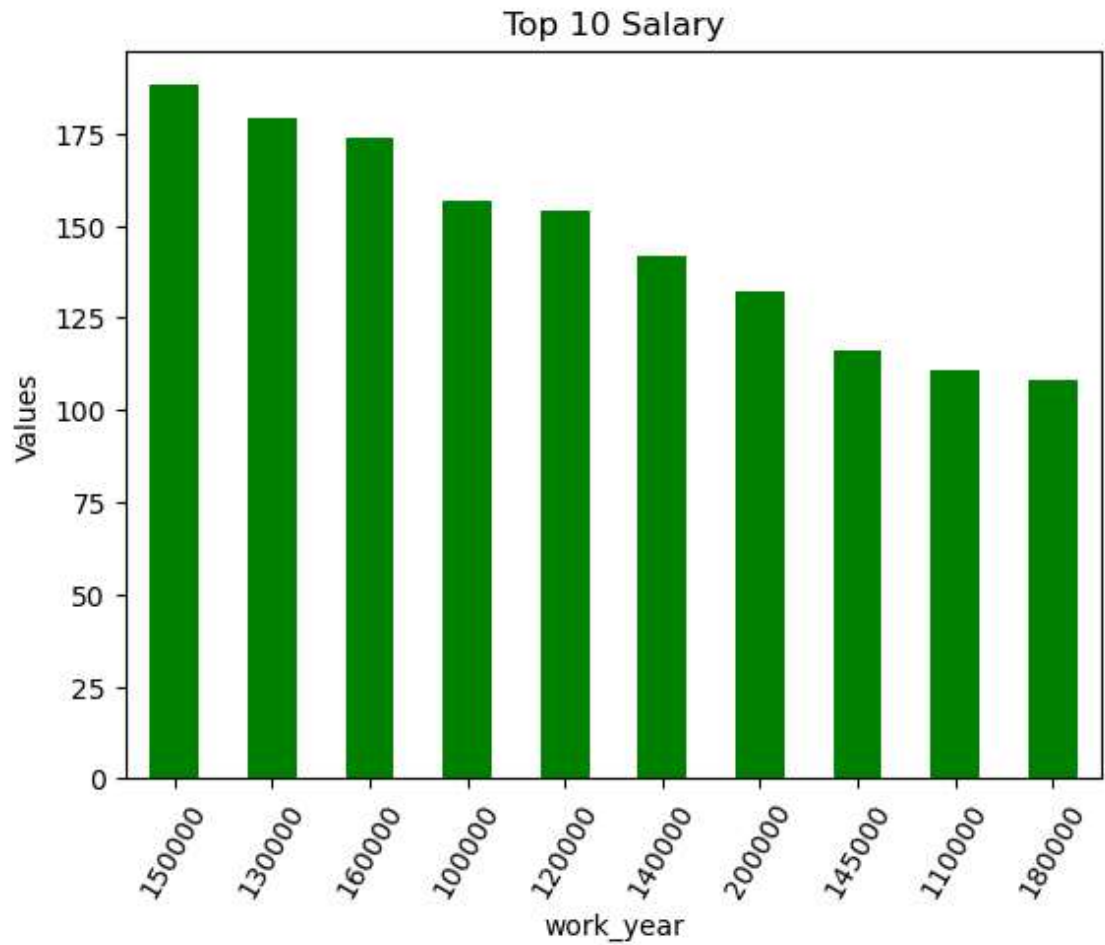
```
Median Salary: 142200.0
Salary Variance: 4124478778.131113
Salary Range: 600201
Mode Salary: [150000]
```

## Building various possible visualizations of dataset variables

```
In [10]: ▶ plt.figure(figsize=(10, 6))  
sns.histplot(df['salary_in_usd'], bins=30, kde=True)  
plt.title('Histogram of Salary in USD')  
plt.show()
```




```
In [20]: Country=df["salary_in_usd"].value_counts()  
Country[:10].plot(kind="bar",color="green")  
plt.title("Top 10 Salary")  
plt.xlabel("work_year")  
plt.ylabel("Values")  
plt.xticks(rotation=60)  
plt.show()
```



```
In [18]: ▶ data=df["employee_residence"].value_counts()
label=df["employee_residence"].unique()

fig,res=plt.subplots()
res.pie(data,labels=label,autopct="%1.1f%%",startangle=45)
plt.show()
```



## Checking any records in the dataset with missing values and handling them appropriately

In [30]: `df.isnull()`

Out[30]:

	work_year	experience_level	employment_type	job_title	salary	salary_currency	salary_in_usd
0	False	False	False	False	False	False	False
1	False	False	False	False	False	False	False
2	False	False	False	False	False	False	False
3	False	False	False	False	False	False	False
4	False	False	False	False	False	False	False
...	...	...	...	...	...	...	...
8800	False	False	False	False	False	False	False
8801	False	False	False	False	False	False	False
8802	False	False	False	False	False	False	False
8803	False	False	False	False	False	False	False
8804	False	False	False	False	False	False	False

8805 rows × 11 columns



In [31]: `df.isnull().sum()`

Out[31]:

work_year	0
experience_level	0
employment_type	0
job_title	0
salary	0
salary_currency	0
salary_in_usd	0
employee_residence	0
remote_ratio	0
company_location	0
company_size	0
dtype:	int64

```
In [32]: df.duplicated()
```

```
Out[32]: 0      False
         1      False
         2      False
         3       True
         4      False
         ...
        8800     False
        8801     False
        8802     False
        8803     False
        8804     False
        Length: 8805, dtype: bool
```


```
In [33]: df.drop_duplicates(inplace=True)
```

```
In [34]: df.duplicated().sum()
```

```
Out[34]: 0
```

## Dealing with Outliers

```
In [21]: ▶ fg = px.box(df,y="salary")  
fg.show()
```




```
In [23]: ▶ Q1 = df.salary.quantile(0.25)  
Q3 = df.salary.quantile(0.75)  
IQR = Q3-Q1
```

```
In [24]: ▶ LowerFence = Q1-1.5*IQR  
UpperFency = Q3+1.5*IQR  
FilteredDF = df[(df.salary>=LowerFence) & (df.salary<=UpperFency)]
```



```
In [25]: ▶ fg = px.box(FilteredDF,y="salary")  
fg.show()
```



## Investigate the unique values of categorical variables


```
In [72]: ▶ categorical_columns = ['experience_level', 'employment_type',  
                                  'salary_currency', 'employee_residence', 'remote_ra  
                                  'company_location', 'company_size']  
unique_values = {col: df[col].unique() for col in categorical_columns}
```

In [73]: `print(unique_values)`

```
{'experience_level': array(['EX', 'MI', 'SE', 'EN'], dtype=object), 'employment_type': array(['FT', 'FL', 'CT', 'PT'], dtype=object), 'salary_currency': array(['USD', 'GBP', 'EUR', 'AUD', 'INR', 'CAD', 'PLN', 'BRL', 'PHP', 'TRY', 'CHF', 'NOK', 'JPY', 'ZAR', 'HKD', 'SGD', 'ILS', 'THB', 'HUF', 'DKK', 'MXN', 'CLP'], dtype=object), 'employee_residence': array(['US', 'GB', 'CA', 'IN', 'UA', 'LT', 'PT', 'NL', 'ES', 'AU', 'FR', 'UG', 'CO', 'IT', 'SI', 'RO', 'GR', 'LV', 'MU', 'DE', 'PL', 'AM', 'HR', 'TH', 'KR', 'EE', 'TR', 'PH', 'BR', 'QA', 'RU', 'KE', 'TN', 'GH', 'BE', 'CH', 'AD', 'EC', 'PE', 'MX', 'MD', 'NG', 'SA', 'NO', 'AR', 'EG', 'UZ', 'GE', 'JP', 'ZA', 'HK', 'CF', 'FI', 'IE', 'IL', 'AT', 'SG', 'SE', 'KW', 'CY', 'BA', 'PK', 'IR', 'AS', 'HU', 'CN', 'CR', 'CL', 'PR', 'DK', 'BO', 'DO', 'ID', 'AE', 'MY', 'HN', 'CZ', 'DZ', 'VN', 'IQ', 'BG', 'JE', 'RS', 'NZ', 'LU', 'MT'], dtype=object), 'remote_ratio': array([ 0, 100, 50], dtype=int64), 'company_location': array(['US', 'GB', 'CA', 'IN', 'NL', 'LT', 'PT', 'GI', 'AU', 'FR', 'CO', 'UA', 'SI', 'RO', 'ES', 'GR', 'LV', 'MU', 'DE', 'PL', 'RU', 'IT', 'KR', 'EE', 'CZ', 'CH', 'BR', 'QA', 'KE', 'DK', 'GH', 'SE', 'PH', 'TR', 'AD', 'EC', 'MX', 'IL', 'NG', 'SA', 'NO', 'AR', 'JP', 'ZA', 'HK', 'CF', 'FI', 'IE', 'SG', 'TH', 'HR', 'AM', 'BA', 'PK', 'IR', 'BS', 'HU', 'AT', 'PR', 'AS', 'BE', 'ID', 'EG', 'AE', 'MY', 'HN', 'DZ', 'IQ', 'CN', 'NZ', 'CL', 'MD', 'LU', 'MT'], dtype=object), 'company_size': array(['M', 'S', 'L'], dtype=object)}
```

## Linear Regression Model

```
In [27]: fig=px.scatter(data_frame=df,x="salary",y="work_year",size="work_year",tree
fig.show())
```



```
In [28]: features=["work_year","salary_in_usd","remote_ratio"]
x=df[features]
y=df["salary"]
```

```
In [29]: xtrain,xtest,ytrain,ytest=train_test_split(x,y,test_size=0.2)
```

```
In [30]: ml=LinearRegression()
ml.fit(xtrain,ytrain)
```

```
Out[30]: LinearRegression
LinearRegression()
```

```
In [31]: ▶ print("Accuracy: ",ml.score(xtest,ytest))
```

Accuracy: 0.004417194198471863